

# ADLR Final Presentation: Development of a reinforcement learning based controller for a VTOL drone.

Christopher Narr & Oliver Hausdörfer

Supervisor: Finn Süberkrüb



TUM / DLR Institute of Robotics and Mechatronics

Advanced Deep Learning for Robotics

23rd February 2023

[c.narr@tum.de](mailto:c.narr@tum.de)

[oliver.hausdoerfer@tum.de](mailto:oliver.hausdoerfer@tum.de)



# Agenda

Topic

Main results

Deep dives

# 1. Topic


**Developing a RL-based velocity-tracking controller for VTOL (Vertical Take Off and Landing) drones**

**Our contributions:**

- + Model time-dependencies
- + Use active wing flaps
- + Incorporate gusty conditions

**Midterm status:**

- Good tracking of simple trajectories
- Using FNN for actor network
- **ToDo:** model time dependencies



# 1. Topic

## Midterm status


RL-Algorithm: PPO

State space (9D): rotation [x,y,z]  
 rotational velocity [x,y,z]  
 linear velocity error [x,y,z]

Action space (4D): 2x thrusts, 2x active flaps

Reward:  $r = 4.0 - 0.5 * ||\mathbf{v}_d - \mathbf{v}_t||_1 - 0.01 * ||\boldsymbol{\omega}_B||_1 - 0.01 * ||\mathbf{a}||_1$

stay alive	tracking error	rotational velocity	actuation
---------------	-------------------	------------------------	-----------




Blue: target velocity | Yellow: actual velocity | Red: thrusts

# 1. Topic

**Midterm status:** Actor Network

FNN (Baseline)

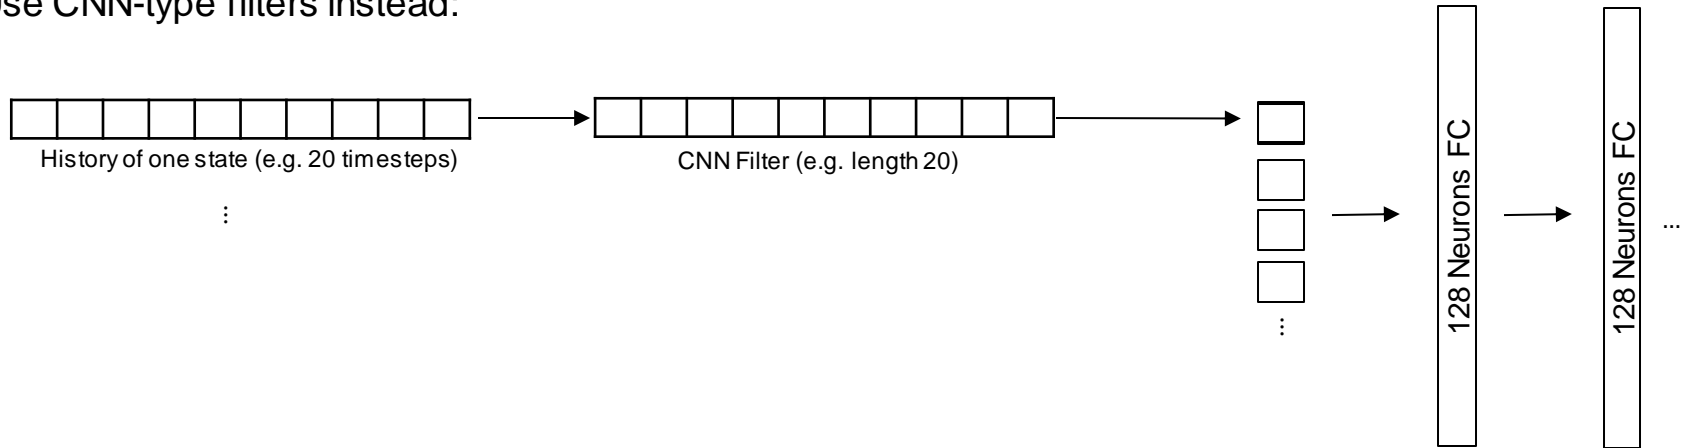


# 1. Topic

**Initial Idea:** Use RNN in actor network for modelling time-dependency

**But:** RNN for RL not supported in Julia “ReinforcementLearning” Library [GitHub Issue #144](#)

Use CNN-type filters instead:



# Agenda

Topic

Main results

Deep dives

## 2. Main results

### Performance of network architectures on test set of 100 envs

Network	Reward	Tracking error (m/s)	# early term. envs	# trainable params
FNN (Baseline)	2696	1.214	4	36.7k
<i>Shared filters (classical CNN)</i>				
1 Filter → FNN	2664	1.231	2	36.7k
3 Filters → FNN	2531	1.586	4	41.1k
5 Filters → FNN	2691	1.158	3	46.1k
9 Filters → FNN	2714	1.138	2	55.5k
18 Filters → FNN	2820	0.899	1	76.6k

### Result

State history significantly improves performance



## 2. Main results

### Performance of network architectures on test set of 100 envs

Network	Reward	Tracking error (m/s)	# early term. envs	# trainable params
FNN (Baseline)	2696	1.214	4	36.7k
<i>Shared filters (classical CNN)</i>				
1 Filter → FNN	2664	1.231	2	36.7k
3 Filters → FNN	2531	1.586	4	41.1k
5 Filters → FNN	2691	1.158	3	46.1k
9 Filters → FNN	2714	1.138	2	55.5k
18 Filters → FNN	2820	0.899	1	76.6k
<i>State dependent filters</i>				
2*9 Filters → FNN	2924	0.62	0	39.1k
4*9 Filters → FNN	2951	0.56	0	45.1k

### Result

State history significantly improves performance

# Agenda


Topic

Main results

Deep dives

# 3. Deep Dive

## Modelling time dependencies: Shared filters (classical CNN) Implementation



Network	Reward	Tracking error (m/s)	# early term. envs	# trainable params
FNN (Baseline)	2696	1.214	4	36.7k
<i>Shared filters (classical CNN)</i>				
1 Filter → FNN	2664	1.231	2	36.7k
3 Filters → FNN	2531	1.586	4	41.1k
5 Filters → FNN	2691	1.158	3	46.1k
9 Filters → FNN	2714	1.138	2	55.5k
18 Filters → FNN	2820	0.899	1	76.6k

# 3. Deep Dive

## Modelling time dependencies: Shared filters (classical CNN)

Analyzing filter length

Network	Reward	Tracking error (m/s)	# early term. envs	# trainable params
FNN (Baseline)	2696	1.214	4	36.7k
<i>Shared filters (classical CNN)</i>				
1 Filter → FNN	2664	1.231	2	36.7k
3 Filters → FNN	2531	1.586	4	41.1k
5 Filters → FNN	2691	1.158	3	46.1k
9 Filters → FNN	2714	1.138	2	55.5k
18 Filters → FNN	2820	0.899	1	76.6k



Activation of weights of trained CNN filter with length 100 (=2.5s)

Reward : 2415 | Tracking Error: 1.7m/s

## Result

Only short-term dependencies relevant

# 3. Deep Dive


## Modelling time dependencies: Shared filters (classical CNN)

Analyzing filter number

Low filter number reduces performance

Network	Reward	Tracking error (m/s)	# early term. envs	# trainable params
FNN (Baseline)	2696	1.214	4	36.7k
<i>Shared filters (classical CNN)</i>				
1 Filter → FNN	2664	1.231	2	36.7k
3 Filters → FNN	2531	1.586	4	41.1k
5 Filters → FNN	2691	1.158	3	46.1k
9 Filters → FNN	2714	1.138	2	55.5k
18 Filters → FNN	2820	0.899	1	76.6k

**Result:** Increase filter number






# 3. Deep Dive

## Modelling time dependencies: Shared filters (classical CNN)

Analyzing filter activation (Take with grain of salt!)

Example feature: rotational velocity around z

Network	Reward	Tracking error (m/s)	# early term. envs	# trainable params
FNN (Baseline)	2696	1.214	4	36.7k
<i>Shared filters (classical CNN)</i>				
1 Filter → FNN	2664	1.231	2	36.7k
3 Filters → FNN	2531	1.586	4	41.1k
5 Filters → FNN	2691	1.158	3	46.1k
9 Filters → FNN	2714	1.138	2	55.5k
18 Filters → FNN	2820	0.899	1	76.6k




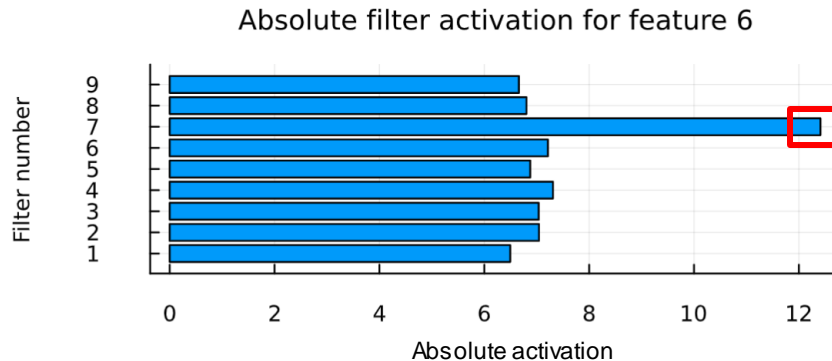
# 3. Deep Dive

## Modelling time dependencies: Shared filters (classical CNN)

Analyzing filter activation (Take with grain of salt!)

Example feature: rotational velocity around z

Network	Reward	Tracking error (m/s)	# early term. envs	# trainable params
FNN (Baseline)	2696	1.214	4	36.7k
<i>Shared filters (classical CNN)</i>				
1 Filter → FNN	2664	1.231	2	36.7k
3 Filters → FNN	2531	1.586	4	41.1k
5 Filters → FNN	2691	1.158	3	46.1k
9 Filters → FNN	2714	1.138	2	55.5k
18 Filters → FNN	2820	0.899	1	76.6k



## Results

Filters are selected

Filters might be state dependent


Filters are 'meaningful'

# 3. Deep Dive

## Modelling time dependencies: State dependent filters

### Implementation

State dependent filters				
2*9 Filters → FNN	2924	0.62	0	39.1k
4*9 Filters → FNN	2951	0.56	0	45.1k





# 3. Deep Dive

## Modelling time dependencies: State dependent filters

Network	Reward	Tracking error (m/s)	# early term. envs	# trainable params
FNN (Baseline)	2696	1.214	4	36.7k
<i>Shared filters (classical CNN)</i>				
1 Filter → FNN	2664	1.231	2	36.7k
3 Filters → FNN	2531	1.586	4	41.1k
5 Filters → FNN	2691	1.158	3	46.1k
9 Filters → FNN	2714	1.138	2	55.5k
18 Filters → FNN	2820	0.899	1	76.6k
<i>State dependent filters</i>				
2*9 Filters → FNN	2924	0.62	0	39.1k
4*9 Filters → FNN	2951	0.56	0	45.1k

## Results

Better performance

Less parameters

# 3. Deep Dive

## Modelling time dependencies: State dependent filters

Network	Reward	Tracking error (m/s)	# early term. envs	# trainable params
FNN (Baseline)	2696	1.214	4	36.7k
<i>Shared filters (classical CNN)</i>				
1 Filter → FNN	2664	1.231	2	36.7k
3 Filters → FNN	2531	1.586	4	41.1k
5 Filters → FNN	2691	1.158	3	46.1k
9 Filters → FNN	2714	1.138	2	55.5k
18 Filters → FNN	2820	0.899	1	76.6k
<i>State dependent filters</i>				
2*9 Filters → FNN	2924	0.62	0	39.1k
4*9 Filters → FNN	2951	0.56	0	45.1k

Total energy consumption on test trajectories:

1.965

1.767

1.701

Less energy consumption indicate more stable flights

## Results

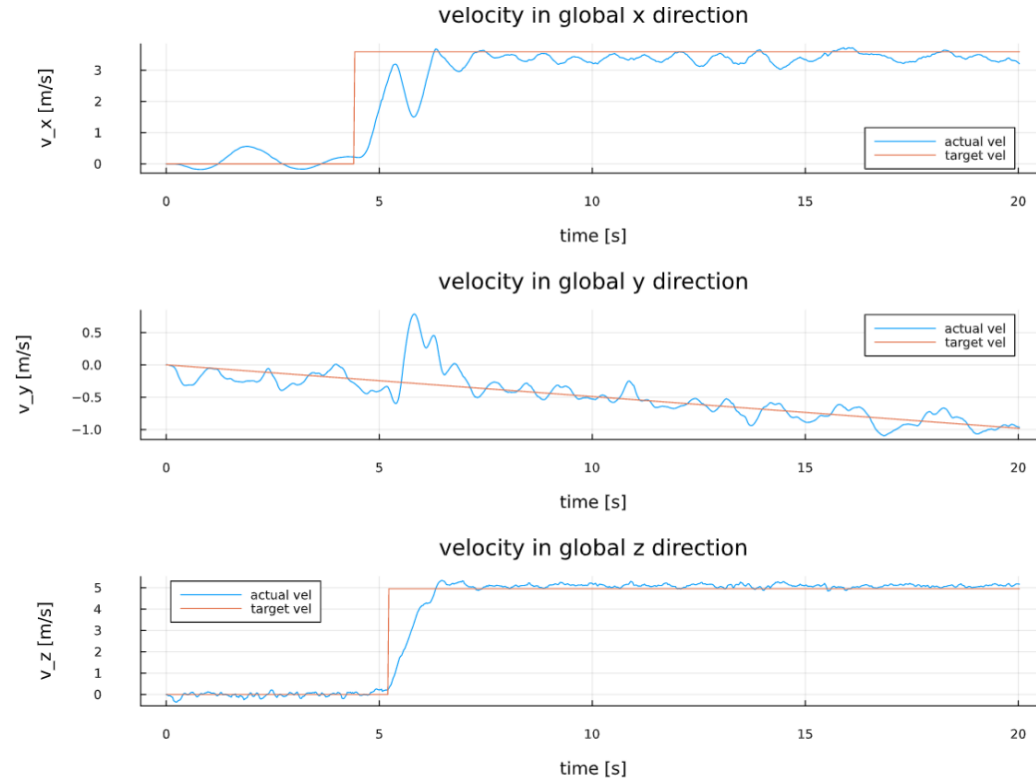
Better performance  
Less parameters  
More stable flights

# 3. Deep Dive

## One example trajectory

Reward: 2897




Tracking error: 0.68 m/s



## Bonus

### Drone to fixed-wing mode

Current velocity as input didn't improve performance



Green: target velocity | Yellow: actual velocity | Red: thrusts

## Bonus

### Other results

- Wind speed as input didn't improve performance
  - Maybe used too low wind speeds (5m/s max)
- Two step training is a double edge sword
  - Updated reward function (less stay alive reward)

Network	Reward	Tracking error (m/s)	# early terminated envs
18 filters → FNN (Baseline)	2820	0.899	1
Add wind speed [x,y,z] as input	2816	0.879	2
Two-step training	-	0.667	13

# Summary & Outlook

- Why has no one done this?
- CNN might be better than RNN for our case
- Good bias improves performance
- State-history control?
- Benchmarks for comparison?

Network	Reward	Tracking error (m/s)	# early term. envs	# trainable params
FNN (Baseline)	2696	1.214	4	36.7k
<i>Shared filters (classical CNN)</i>				
1 Filter → FNN	2664	1.231	2	36.7k
3 Filters → FNN	2531	1.586	4	41.1k
5 Filters → FNN	2691	1.158	3	46.1k
9 Filters → FNN	2714	1.138	2	55.5k
18 Filters → FNN	2820	0.899	1	76.6k
<i>State dependent filters</i>				
2*9 Filters → FNN	2924	0.62	0	39.1k
4*9 Filters → FNN	2951	0.56	0	45.1k

# ADLR Final Presentation: Development of a reinforcement learning based controller for a VTOL drone.

Christopher Narr & Oliver Hausdörfer

Supervisor: Finn Süberkrüb

TUM / DLR Institute of Robotics and Mechatronics

Advanced Deep Learning for Robotics

23rd February 2023

Thank you!

Q&A

